

From a 10-page spec to a shippable MVP.

The Dashboard Content Guide defines a large analytics surface — ten pages, dozens of charts, basic and advanced modes, prop, live, and multi-source views. This plan sequences that scope into three phases so the first usable, trustworthy product ships at the end of Phase 1.

PHASE 1 · MVP

The Trust Decision Layer

One source, one Overview page.
Answers all four core questions.
Scalar data only.

PHASE 2

Depth, Filters & Compare

Filtering, comparison, prop, and the
analysis pages that need richer
fields.

PHASE 3

Advanced & Multi-Source

Price-series analytics, simulation, live
discipline, source reconciliation.

SOURCE `DASHBOARD_CONTENT_GUIDE.md`

IMPL `src/TalariaV8b.jsx`

BASIS 4 core trader questions

Three gates set every feature's phase

Each item from the content guide is scored on three independent gates. A feature lands in the *latest* phase that any gate forces — you cannot ship something before its data exists, before it earns its place, or before it is feasible to build well.

Gate A — Data Dependency What raw fields must exist for this visual to be correct? A chart cannot ship before its inputs do. Tier A → Phase 1 Tier B → Phase 2 Tier C → Phase 3	Gate B — Decision Value How directly does it answer the four core questions: profitable, stable, risk acceptable, rules followed? Core answer → Phase 1 Supporting depth → Phase 2 Exploratory → Phase 3	Gate C — Build Complexity Effort and risk: simple aggregation, a stateful engine, or heavy simulation / cross-source plumbing? Low (aggregate) → ok early Medium (engine) → Phase 2 High (sim / sync) → Phase 3
---	--	---

THE RULE Phase = the maximum across all three gates. This is why a few features whose data is trivially available (Runs Test, Edge Decay) still land in Phase 3 — their value is exploratory and their build is heavy, so a later gate dominates. It also keeps Phase 1 honest: nothing enters the MVP unless it is core, cheap, *and* buildable on data that always exists.

GATE A IN DETAIL · THE DATA TIERS

the dominant constraint

Tier A Always present · scalar · no series	Per trade: date, symbol, market, side, pnl, rMultiple, duration, tag, outcome, rulesFollowed, rulesViolated. Per session: capital, dates, tradingMode. Unlocks: every KPI, equity & drawdown curves, win/loss, R-distribution, daily heatmap, breakdowns, the trade table, basic rule adherence, no-go flags, a static Talaria Score.
Tier B Enriched · needs tagging or excursion capture	Per trade: preTags, postTags, plannedRR, actualRR, entryQuality, exitQuality, mae, mfe, positionSize, marketRegime, nearbyNewsEvent. Plus prop preset + prior-period metrics. Unlocks: MAE/MFE scatter, planned-vs-actual execution, behavioral tags, edge finder, prop challenge, quant risk strip, goal tracker, full Talaria Score radar & trend.
Tier C Heavy · price-series, simulation, cross-source	Per trade: bar_close_r[] / bar_high_r[] / bar_low_r[] price arrays. Cross-source: strategy→session/journal hierarchy, live sync, edit flags, symbol return series. Unlocks: Price Behavior Explorer, Monte Carlo, correlation matrix, what-if intra-bar replay, live discipline, source reconciliation, report builder.

PHASE 1

The Trust Decision Layer

A trader loads one backtest or one journal and learns — on a single Overview page — whether the source is profitable, stable, risk-acceptable, and rule-followed. Nothing here needs tags, excursion, prop presets, price series, or a second source.

■ MVP
MILESTONE
Tier A data only

IN SCOPE — KPI STRIP & VERDICT

- ◆ Net P&L, Return %
did it make money, and is it meaningful vs account size
- ◆ Trade Win Rate · Day Win Rate
how often it wins, and whether wins are steady or concentrated
- ◆ Profit Factor · Avg R
structurally profitable, and edge survives size normalization
- ◆ Max Drawdown · Trades / Duration
worst pain, plus sample-size sanity
- ◆ Talaria Score
single static composite (no trend/radar yet)
- ◆ No-Go Flags
drawdown, thin sample, rule breaks, negative expectancy

IN SCOPE — CORE CHARTS

- ◆ Equity Curve & Drawdown Curve
- ◆ Win / Loss donut
- ◆ R-Multiple Distribution
- ◆ Daily P&L Heatmap
- ◆ Performance Breakdown
by symbol / tag / outcome
- ◆ Trade Journal table
row-level review of filtered trades
- ◆ Rule Adherence panel
followed vs broke: P&L, R, win rate

IN SCOPE — FOUNDATIONS

- ◆ Source button
single-source selection + the core data model
- ◆ Value mode
\$ / % / privacy — privacy must hide every dollar from day one
- ◆ Page nav shell
scaffolding only; Overview is the one live page
- ◆ Empty & edge states
0 trades, 1 trade, 1000+ trades must never crash
- ◆ RTL / Arabic label support
readable, unclipped — non-negotiable for MENA

RECOMMENDED BUILD ORDER

- 01 Data layer + Source model — normalize Source / Session / Trade; wire single-source select.
- 02 Value mode + edge states — build privacy & empty-state handling before any visual.
- 03 KPI strip — the eight cards; pure aggregation, fastest value.
- 04 Equity + Drawdown — the stability story.
- 05 Win/Loss, R-dist, Heatmap, Breakdown — the shape of the sample.
- 06 Journal table + Rule Adherence — the “which trades” and “following rules” answers.
- 07 Talaria Score + No-Go Flags — the at-a-glance verdict; ship.

PHASE 1 DEFINITION OF DONE

- ✓ One backtest *or* one journal loads end-to-end and renders the full Overview page.
- ✓ All four core questions are answerable without leaving that page.
- ✓ Privacy mode hides every dollar value while keeping %, counts, ratios, score, and chart structure.
- ✓ 0 / 1 / 1000+ trades all render without errors; filtered-to-empty shows shells, never a crash.
- ✓ Zero dependency on tags, MAE/MFE, prop presets, price arrays, or compare mode.

EXPLICITLY NOT IN PHASE 1

Deferred: filters, compare mode, Add/Edit Trade, prop banner, in-progress panel, edge finder, goal tracker, score trend/radar, and every page other than Overview. Pushing these out is what makes a Phase 1 ship date real.

PHASE 2

Depth, Filtering & Comparison

Turn the MVP into a real analysis tool: slice the data, compare two sources, and open the pages that read tags, execution quality, and prop pressure. This is where Tier B data and the filter engine arrive.

■ ANALYSIS DEPTH
Tier A + B data

ENGINE & GLOBAL CONTROLS

- ◆ Filters button
date, timing, trade scope, pre/post tags — drives all charts
- ◆ Compare button + summary banner + Comparison View
- ◆ Add / Edit Trade flow
with source edit-status flags
- ◆ Source status icons
rollback, costs, edited/original
- ◆ In-progress backtest panel

NEW PAGES

- ◆ Performance
equity & returns, drawdown, consistency, risk-adjusted
- ◆ The Numbers
distributions, calendar & time, breakdowns, cross-analysis
- ◆ Trade Quality
MAE/MFE scatter, planned vs actual, management quality
- ◆ Prop Challenge
rule table, pass/fail pressure, status banner
- ◆ Journal (full)
filtered summary + rich row review

NEW VISUALS & INSIGHT LAYERS

- ◆ Quant KPI Strip
Sharpe, Sortino, Calmar, Recovery Factor
- ◆ Edge Finder strip
plain-language strengths & weaknesses
- ◆ Talaria Score — full
six-dimension radar + score trend
- ◆ Performance Calendar · Consistency · Streaks & Records
- ◆ Behavioral Tags
performance by pre/post tag
- ◆ Goal Tracker
- ◆ Rule Compliance Tracker
prop rule nearest to failing

PHASE 2 DEFINITION OF DONE

- ✓ Every chart re-computes correctly on source change *and* filter change.
- ✓ Compare mode reads the right main vs compare source on every metric.
- ✓ Prop rules and the prop banner appear only for prop-mode sources.
- ✓ Tag-, execution-, and risk-adjusted views render from Tier B fields.

STILL NOT IN PHASE 2

Deferred to Phase 3: anything needing per-bar price arrays, Monte Carlo or other simulation, correlation series, live-journal sync, or multi-source reconciliation. Phase 2 stays inside data the trade record already holds.

Advanced Analytics & Multi-Source

PHASE 3

The differentiation layer: heavy simulation, price-path analytics, behavioral edge work, live-execution discipline, and the reconciliation that ties backtest, prop, live, and strategy sources into one picture. Requires Tier C data and infrastructure.

DIFFERENTIATION
Tier C + cross-source

PRICE-SERIES & SIMULATION

- ◆ Price Behavior Explorer
avg price paths around entries
- ◆ Monte Carlo Projection & Pass Probability
- ◆ What-If Lab
simulator + alternative strategies (intra-bar replay)
- ◆ Correlation Matrix · 3D Heatmap
- ◆ Advanced Distributions

SEQUENCE & BEHAVIOR

- ◆ Edge Decay
rolling expectancy
- ◆ Sequence Risk · Runs Test
- ◆ Trade Microscope
by duration, volatility, regime
- ◆ Macro Event Correlation
trades near news
- ◆ Edge & Behavior page
patterns, sabotage points, strategy health

LIVE & MULTI-SOURCE

- ◆ Live Discipline page
live plan adherence, edited-journal status
- ◆ Strategy & Sources page
reconciliation, attribution, backtest vs live
- ◆ Multiple-source aggregation with attribution preserved
- ◆ Report Builder
section selection & export

PHASE 3 DEFINITION OF DONE

- ✓ Price-path and excursion visuals render accurately from per-bar arrays.
- ✓ Simulation views (Monte Carlo, what-if) reflect real intra-bar outcomes.
- ✓ Live journal vs backtest reconcile under one strategy with attribution intact.
- ✓ Reports export the selected pages/sections faithfully.

SEQUENCING NOTE

Phase 3 items are largely independent — ship them by audience demand. **Prop-heavy users** → Monte Carlo / pass probability first. **Live traders** → Live Discipline first. **Quants** → price-path & what-if first.

Where each feature lands — and why

Phase is the maximum of the three gates. T = data tier required. The reason column names the gate that pushed the item to its phase.

FEATURE	PHASE	T	DRIVING REASON
FOUNDATIONS & GLOBAL CONTROLS			
Source selection + data model	P1	A	Everything depends on it; foundational.
Value mode (privacy / \$ / %)	P1	A	Cheap, core trust feature; MENA privacy need.
Empty / edge states	P1	A	Must be correct from the first render.
Filters (date, timing, scope, tags)	P2	B	Tag filters need Tier B; stateful engine.
Compare mode + Comparison View	P2	A	Supporting depth, not MVP (Gate B).
Add / Edit Trade	P2	A	MVP is read-only analysis.
Source status icons	P2	B	Needs cost/edit/rollback flags.
OVERVIEW & KPIS			
8 core KPIs (P&L → duration)	P1	A	Direct answers to the core questions.
Rule Adherence panel	P1	A	Core question #4; Tier A fields.
No-Go Flags	P1	A	High decision value, simple thresholds.
Talaria Score — static composite	P1	A	At-a-glance verdict from Tier A inputs.
Talaria Score — radar / trend	P2	B	Needs dimensions + prior-period data.
Edge Finder strip	P2	A	Insight engine; medium complexity (Gate C).
Prop status banner	P2	B	Prop preset data; prop sources only.
In-progress backtest panel	P2	A	Useful, not part of the verdict.
CHARTS			
Equity / Drawdown / Win-Loss / R-dist	P1	A	Profitability + stability core.
Daily Heatmap · Breakdown · Journal table	P1	A	Cheap aggregation, high value.
Quant KPI Strip (Sharpe/Sortino/Calmar)	P2	A	Risk depth beyond MVP (Gate B/C).
Cross-Analysis · Calendar · Consistency	P2	A	Depend on the filter engine.
Behavioral Tags · Goal Tracker	P2	B	Need tags / target inputs.
MAE/MFE Scatter · Execution Analysis	P2	B	Need excursion + planned/actual RR.
Monte Carlo · Price Behavior · Correlation	P3	C	Price series / simulation (Gate A + C).
3D Heatmap · Advanced Distributions	P3	C	Heavy build; exploratory value.
Edge Decay · Sequence Risk · Runs Test	P3	A	Data is cheap, but value+build push late.
Trade Microscope · Macro Event Corr.	P3	B	Regime / news; exploratory.
PAGES			
Overview	P1	A	The MVP page.
Performance · The Numbers · Trade Quality	P2	B	Depth pages; some need Tier B.
Prop Challenge · Journal (full)	P2	B	Prop presets / rich review.
What-If Lab · Edge & Behavior	P3	C	Simulation + behavioral engine.
Live Discipline · Strategy & Sources	P3	C	Live sync + reconciliation infra.
Report Builder / export	P3	C	Depends on pages existing first.

What protects the MVP date

The biggest threat to a Phase 1 ship is scope leaking backward from later phases. These are the dependencies to lock and the traps to avoid.

RISK 1**Building visuals before the data layer.**

If charts are wired to ad-hoc shapes, every Phase 2 field (tags, MAE/ MFE) forces a rewrite. **Mitigation:** normalize the Source / Session / Trade model first, with Tier B/C fields declared as optional from day one — charts read them defensively.

RISK 2**Privacy & empty states bolted on late.**

Retrofitting privacy mode and 0/1-trade handling across finished charts is expensive and bug-prone. **Mitigation:** build both as cross-cutting primitives in step 2, before any visual exists.

RISK 3**The Talaria Score scope-creeping into Phase 1.**

The radar and score-trend need six scored dimensions and prior-period history — a Phase 2 concern. **Mitigation:** ship a single static composite in Phase 1; expand the same score object later without breaking the card.

RISK 4**Filters sneaking into the MVP.**

A full date/timing/scope/tag engine is the heart of Phase 2 and easily doubles Phase 1. **Mitigation:** Phase 1 shows one source, unfiltered. Resist “just a date range.”

RISK 5**RTL discovered at the end.**

Arabic-native is the product's reason to exist; finding clipped/reversed layouts after the fact is costly. **Mitigation:** test RTL on the KPI strip and first chart, not at sign-off.

One-line summary. Phase 1 ships a single-source, unfiltered, privacy-aware Overview built only on always-present scalar fields — the smallest thing that answers *profitable? stable? risky? disciplined?* Phase 2 adds filtering, comparison, prop, and the Tier B analysis pages. Phase 3 adds price-series analytics, simulation, live discipline, and multi-source reconciliation. Keep the data model forward-compatible and the phase boundaries hold.